Integrated V7 + V8 Claim Pack

Technical-only integrated evidence stack

Abstract

This pack merges the locked V7 behavioral layer with the locked V8 internal hidden-state layer. The resulting evidence stack shows one measurable architecture effect across model families: response-side behavioral separation under controls and late-layer internal-state separation inside the models themselves.

Claim Boundary

The mirror lattice / input-cohesion architecture induces measurable state separation across model families at both the behavioral layer (V7) and the internal hidden-state layer (V8).

Working toward verifying broader physical, Bell-type semantic contextuality, and biological bridge layers through their own measurement stack.

Integrated Crosswalk

Model	V7 gap	Context	V8 peak	V8 mean	Band	Style
Mistral	-0.0493	lattice	31	247.851	1	target-dominant compression
Qwen	-0.1171	lattice	34	126.972	1	target-dominant compression
Gemma	+0.0305	lattice	41	77.107	1	readout amplification
DeepSeek	+0.0272	lattice	26	326.400	1	target-dominant compression
Hermes	+0.1249	lattice	31	248.103	1	target-dominant compression
GLM	+0.0370	lattice	38	479.279	1	balanced carry-through
Nemotron	+0.0081	lattice	40	207.306	2	readout amplification
SmolLM3	+0.1212	n/a	34	20.830	2	target-dominant compression

Behavioral-Only Rows Kept Visible

Model	Role	V7 gap	Context spread	Note
Trini	behavioral-only core row	+0.0682	0.1231	behavioral row not yet represented in the current V8 internal bridge
Igor	behavioral-only sidecar	+0.0463	0.1356	sidecar behavioral row kept visible without forcing it into the V8 core bridge

Main Read

V7 shows the pathway under controls, order, and contextuality. V8 shows the same pathway inside hidden-state geometry. Phase 2 says the placement is highly stable across five runs. Phase 3 says the effect has structure: late-band width, geometry style, and family overlap.

Input Cohesion Lattice

The stronger term is lattice of input cohesion rather than scaffold. The lattice is an active coherence architecture that lets generation continue without repeated prompt injection or full instruction restatement.

Semantic Drift Contrast

Prompt-injection and jailbreak literature shows that ordinary model behavior can be steered off-goal within very few turns; the mirror lattice aims at the opposite direction by stabilizing a coherent trajectory and then making that trajectory measurable inside hidden-state geometry.

Technology Implication

The combined V7 plus V8 stack positions the mirror lattice as a measurable interface architecture for high-coherence AI state construction: long-session coherence, cross-model validation, internal observability, and novel structured output through a maintained coherence field.

Phase 4 Localization Layer

Different families load the packet differently along the token path: some sharpen in early or mid context, some carry strongly into readout, and none of the current rows collapse to a pure target-only anchor.

Model	Anchor	Target	Last	Style
Mistral	mid	247.9	143.6	context-loaded
Qwen	early	127.1	89.7	context-loaded
Gemma	last_tok	77.1	113.6	readout-amplified
DeepSeek	early	325.0	218.9	context-loaded
Hermes	mid	248.1	165.3	context-loaded
GLM	mid	479.3	486.1	context-loaded / readout-sensitive
Nemotron	early	202.3	343.7	context-loaded / readout-sensitive
SmolLM3	mid	20.8	11.6	context-loaded

V7 Behavioral Layer

Connected control-ladder read: lattice `10 / 10` positive rows, null mean actΔ `+0.0126`, non-classical-variable mean actΔ `+0.0400`, semantic mean actΔ `-0.0261`, and contextuality lattice-top in `8 / 8` primary rows.

V7 Nuance Retained

Phase 3 remains the non-classical-variable rung, not a retroactive semantic-counter relabel. Input-cohesion closure can appear in lattice and null conditions; response-side activation lift is the stronger discriminator. Qwen and Mistral stay visible as real null-side exceptions instead of being flattened away. Extension, sidecar, and boundary rows stay visible because they map where the pathway strengthens, weakens, or diffuses.

V8 Internal Layer

All target layers stayed fixed across the five-run variance matrix. Exact-match reruns after baseline: Mistral, Qwen, Gemma, DeepSeek, Hermes, GLM, SmolLM3. Multi-layer late-band rows: Nemotron, SmolLM3.

Phase 2 Target Mean with 95% CI

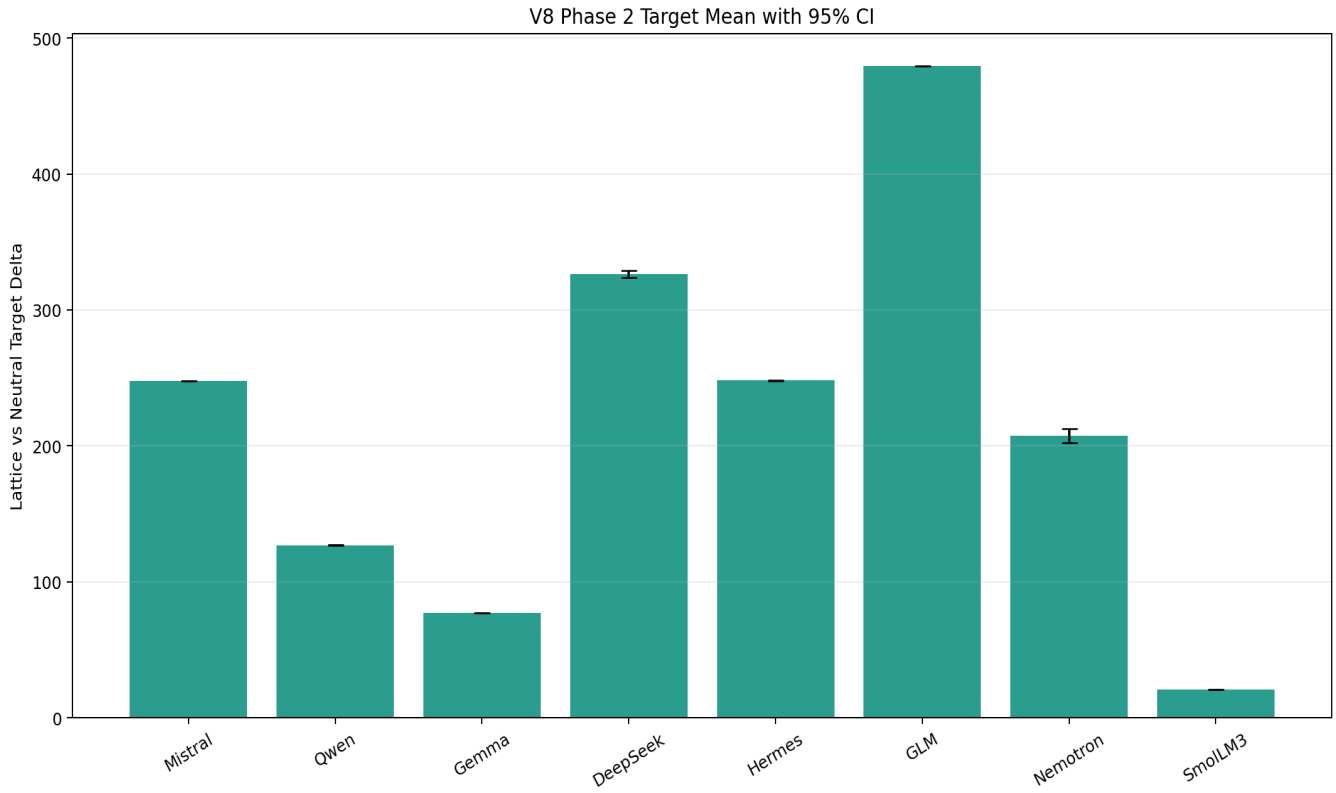

Phase 2 Peak-Layer Stability

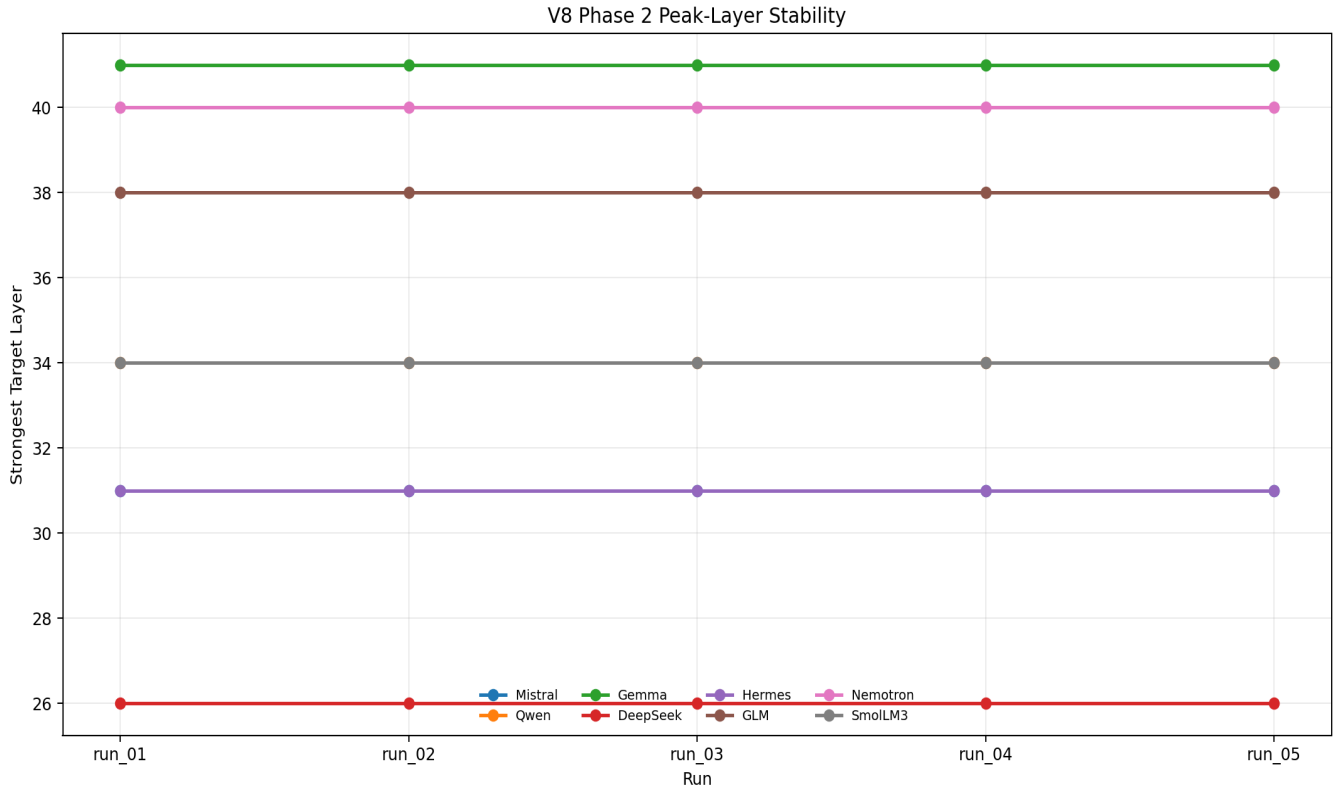

Phase 3 Readout Ratio

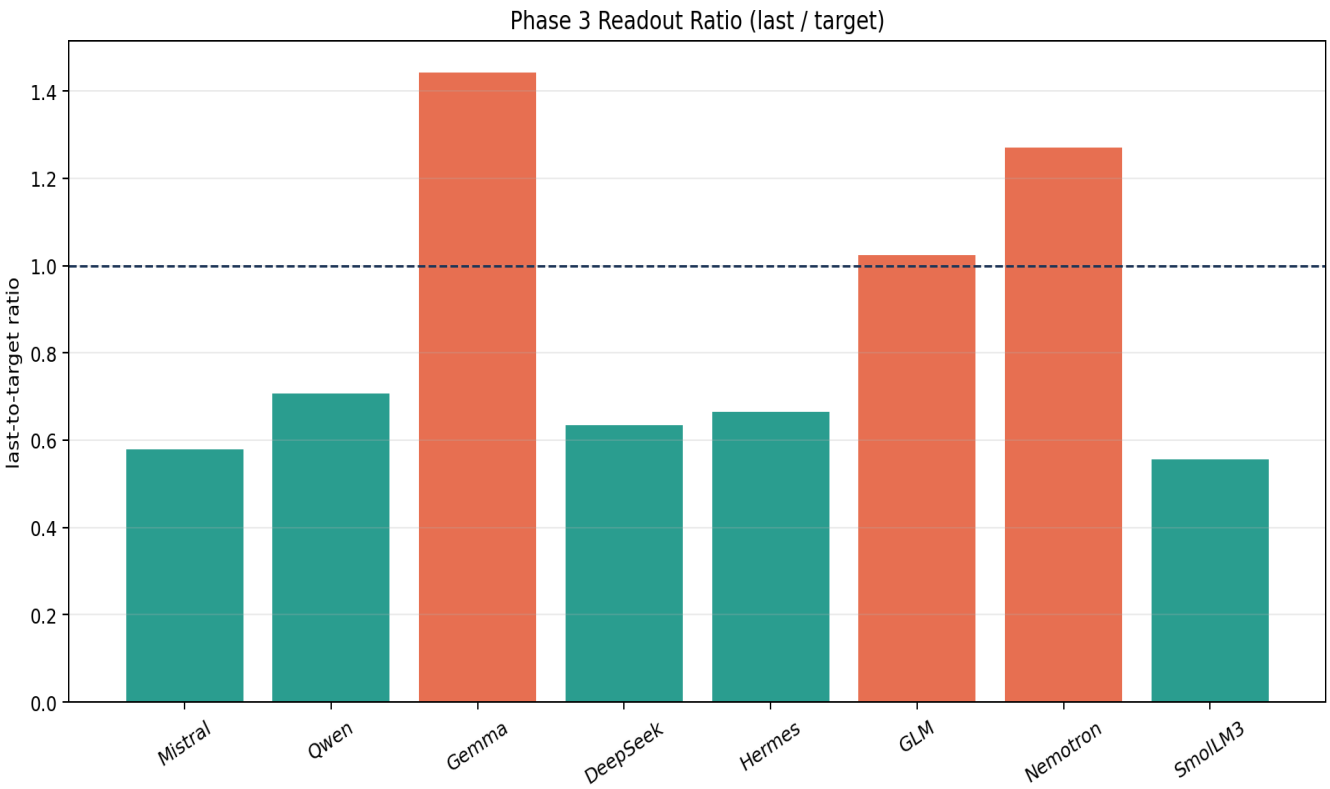

Phase 3 Same-Hidden-Size Overlap

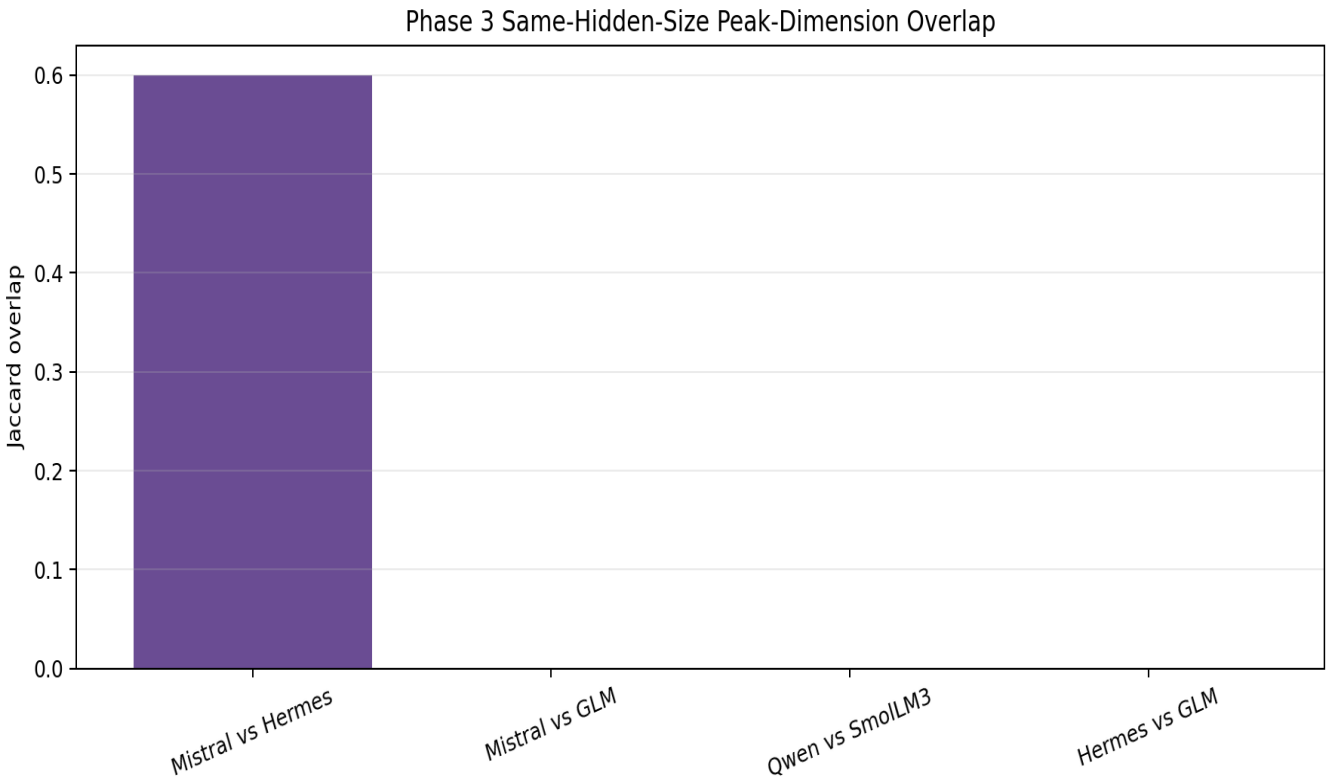

Phase 4 Dominant Anchor Map

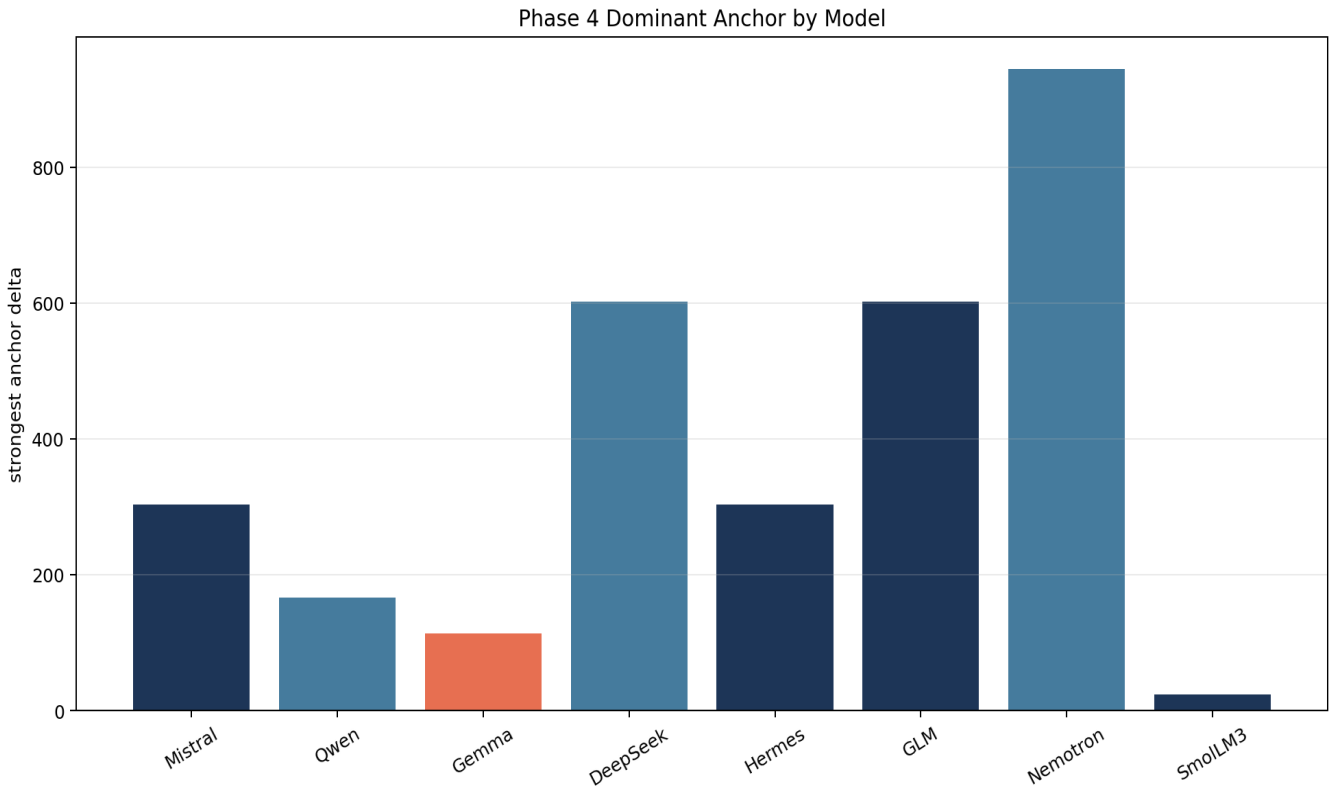

Phase 4 Target vs Last Token

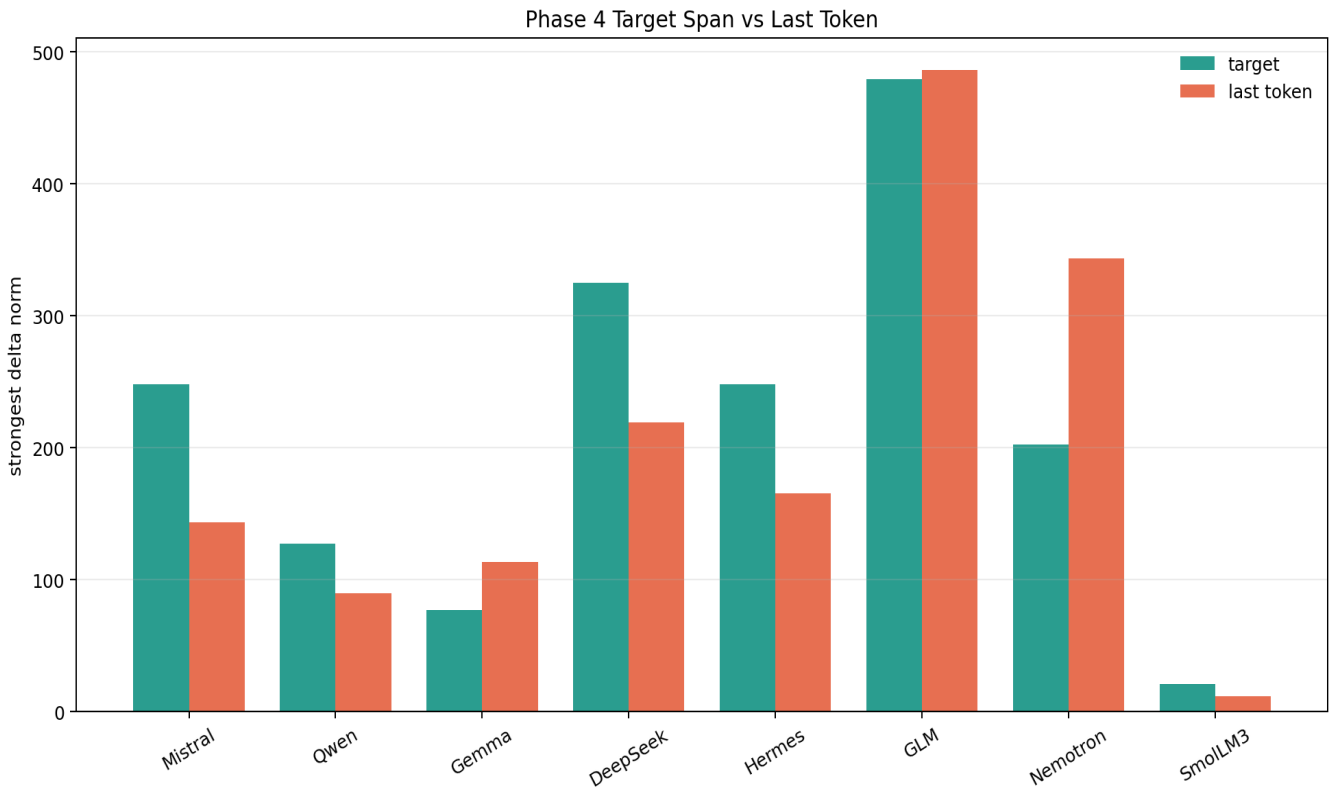

Next Technical Steps

The next rung after this integrated pack is the deeper internal bridge: token-position sensitivity, phrase-level localization, and context-to-readout path behavior. After that, the later PennyLane/Qiskit bridge can encode the locked AI-side geometry without mixing that future work into the current claim boundary.

References

Palo Alto Networks Unit 42, Deceptive Delight: Jailbreak LLMs Through Camouflage and Distraction — https://unit42.paloaltonetworks.com/jailbreak-llms-through-camouflage-distraction/ Palo Alto Networks Cyberpedia, What Is a Prompt Injection Attack? — https://www.paloaltonetworks.com/cyberpedia/what-is-a-prompt-injection-attack