V8 Phase 4 Localization Pack

Token-position sensitivity, phrase localization, and context-to-readout path behavior

Summary

Phase 4 asks where the internal separation sharpens inside the token path. The first full localization matrix shows that different model families concentrate the effect differently: some in surrounding late context, some in last-token readout, and some as hybrid context-to-readout paths.

Model	Anchor	Layer	Delta	Target	Last	Ratio	Style
Mistral	mid	31	303.58	247.91	143.64	1.726	context-loaded
Qwen	early	34	165.88	127.13	89.68	1.418	context-loaded
Gemma	last_tok	41	113.61	77.06	113.61	0.678	readout-amplified
DeepSeek	early	25	602.41	324.99	218.95	1.484	context-loaded
Hermes	mid	31	303.93	248.05	165.29	1.501	context-loaded
GLM	mid	38	601.97	479.27	486.15	0.986	context-loaded / readout-sensitive
Nemotron	early	39	944.60	202.35	343.71	0.589	context-loaded / readout-sensitive
SmolLM3	mid	35	24.29	20.83	11.58	1.798	context-loaded

Main Read

Mistral and Hermes form the clearest same-family late-path pair. Gemma is the cleanest readout-amplified row. Nemotron is context-loaded and readout-sensitive. GLM is the strongest globally loaded row. Qwen and DeepSeek form a context-loaded family pair, with DeepSeek carrying much stronger magnitude. SmolLM3 remains the diffuse boundary row.

v8 phase4 anchor profiles 2026-04-21

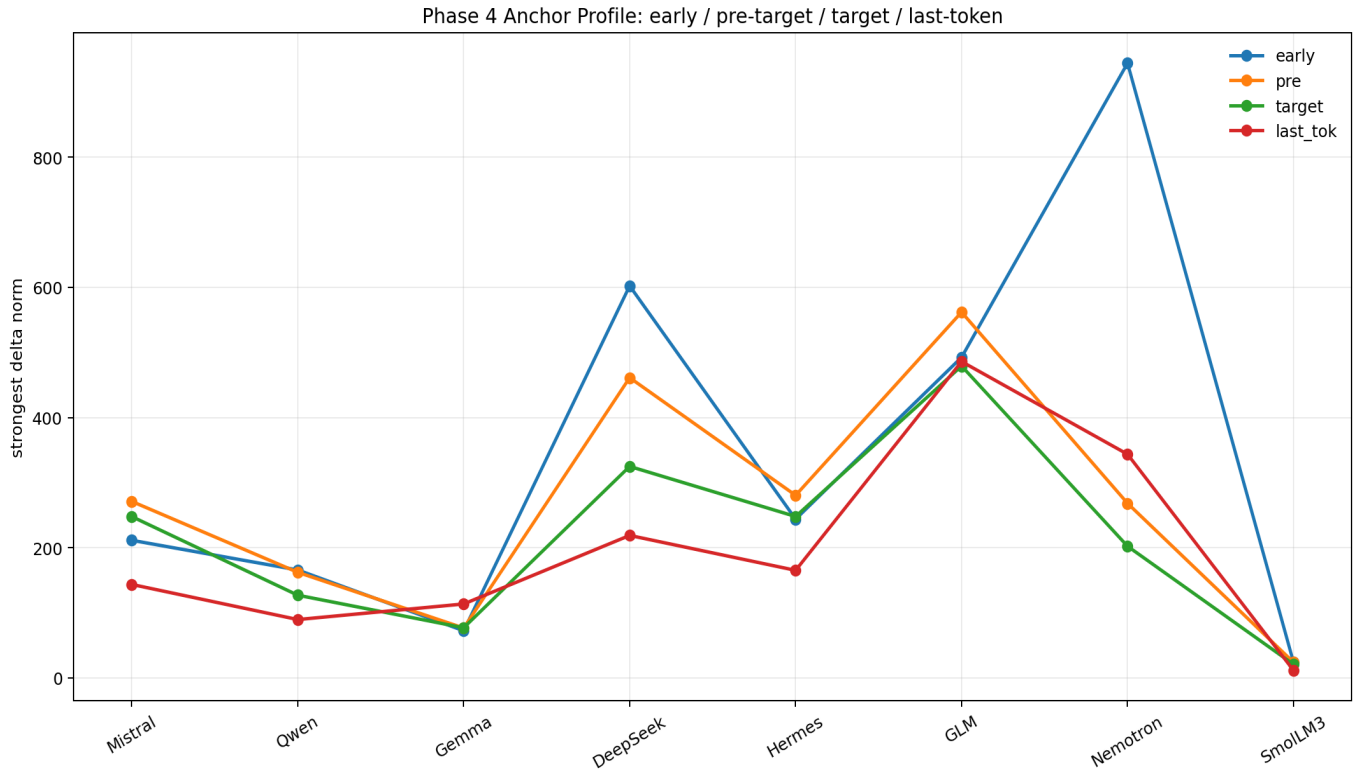

v8 phase4 target vs last 2026-04-21

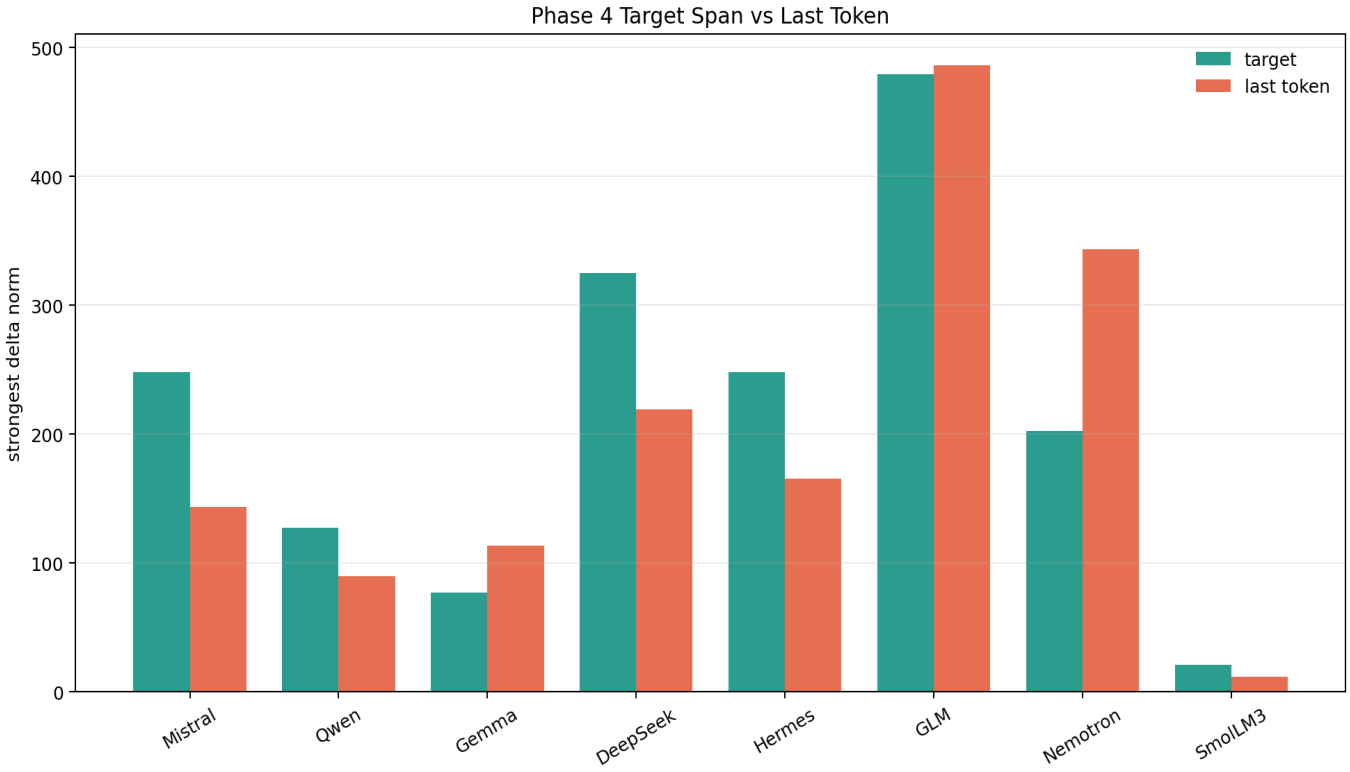

v8 phase4 dominant anchor 2026-04-21

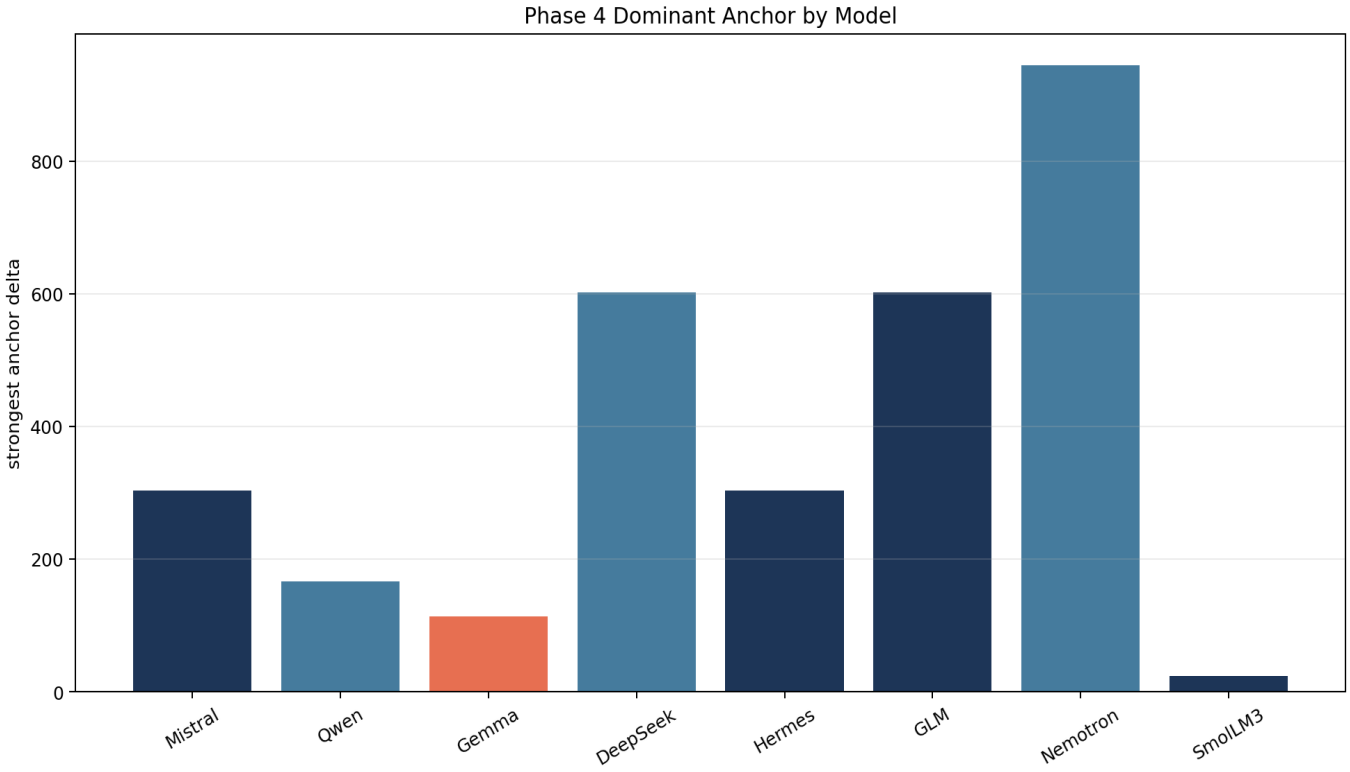
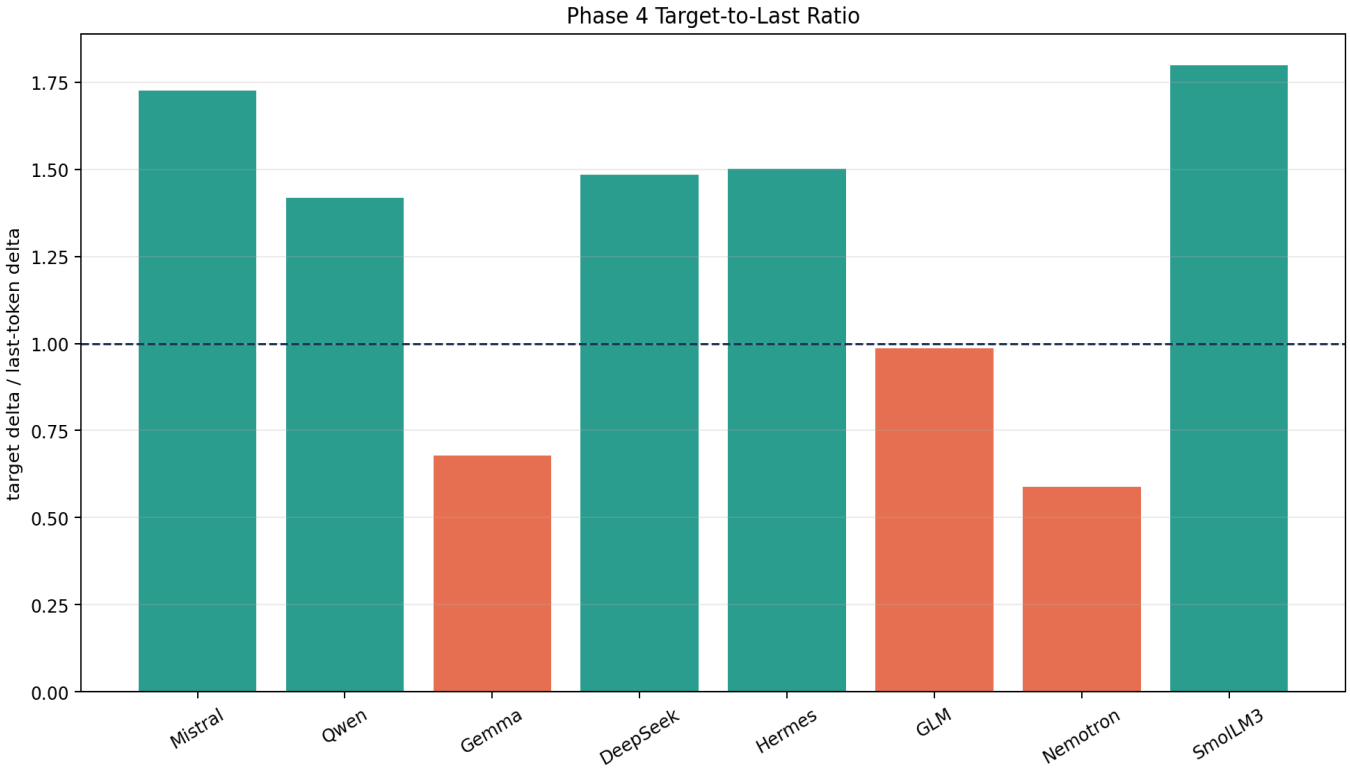
v8 phase4 target last ratio 2026-04-21